

UCI Adult Income Dataset - Exploratory and Descriptive Analysis

In this notebook, we carry out an in-depth exploratory and descriptive analysis of the UCI Adult Income Dataset, a widely used dataset for income prediction tasks based on individual demographic and employment attributes.

This phase of analysis is essential for uncovering patterns, detecting potential biases, and gaining intuition about the dataset's structure before applying any modelling procedures. We examine the distribution of key numerical and categorical variables, investigate relationships between demographic features and income levels, and use visualizations to summarize insights. Particular focus is placed on income disparities across **age groups**, **geographical regions**, **rac**es, and **education-occupation combinations**, helping lay a solid foundation for downstream modeling and policy-relevant interpretation.

We begin our analysis by importing the core Python libraries required for **data handling**, **numerical computation**, **visualization**, and **directory management**:

- **pandas**: Enables efficient manipulation, filtering, and aggregation of structured tabular data, forming the backbone of our analysis pipeline.
- **numpy**: Provides support for fast numerical operations, array-based computation, and statistical routines.
- **os**: Facilitates interaction with the file system, allowing us to construct flexible and portable directory paths for data and output management.
- **plotly.express**: A high-level graphing library that enables the creation of interactive, publication-quality visualizations, which we use extensively to uncover patterns and present insights throughout the notebook.

```
# Import libraries  
  
import pandas as pd  
import numpy as np
```

```
import os
import plotly.express as px
```

Define and Create Directory Paths

To ensure reproducibility and organized storage, we programmatically create directories if they don't already exist for:

- **raw data**
- **processed data**
- **results**
- **documentation**

These directories will store intermediate and final outputs for reproducibility.

```
# Get working directory
current_dir = os.getcwd()
# Go one directory up to the root directory
project_root_dir = os.path.dirname(current_dir)
# define paths to the data files
data_dir = os.path.join(project_root_dir, 'data')
raw_dir = os.path.join(data_dir, 'raw')
processed_dir = os.path.join(data_dir, 'processed')
# Define paths to the results folder
results_dir = os.path.join(project_root_dir, 'results')
# Define paths to the docs folder
docs_dir = os.path.join(project_root_dir, 'docs')

# create directories if they do not exist
os.makedirs(raw_dir, exist_ok = True )
os.makedirs(processed_dir, exist_ok = True )
os.makedirs(results_dir, exist_ok = True)
os.makedirs(docs_dir, exist_ok = True)
```

Loading the Cleaned Dataset

We load the cleaned version of the UCI Adult Income Dataset from the processed data directory into a Pandas DataFrame. The `head(10)` function shows the first ten records, giving a glimpse into the data columns such as `age`, `workclass`, `education_num`, etc

```
adult_data_filename = os.path.join(processed_dir, "adult_cleaned.csv")
adult_df = pd.read_csv(adult_data_filename)
adult_df.head(10)
```

	age	workclass	fnlwgt	education_num	marital_status	relationship	race	sex	ca
0	39	government	77516	13	single	not-in-family	white	male	2
1	50	self-employed	83311	13	married	husband	white	male	0
2	38	private	215646	9	divorced or separated	not-in-family	white	male	0
3	53	private	234721	7	married	husband	black	male	0
4	28	private	338409	13	married	wife	black	female	0
5	37	private	284582	14	married	wife	white	female	0
6	49	private	160187	5	divorced or separated	not-in-family	black	female	0
7	52	self-employed	209642	9	married	husband	white	male	0
8	31	private	45781	14	single	not-in-family	white	female	14
9	42	private	159449	13	married	husband	white	male	5

Dataset Dimensions and Data Types

Here, we examine the structure of the dataset:

- There are *32,513* entries and *16* variables.
- The dataset includes both **numerical** (e.g., `age`, `hours_per_week`) and **categorical** variables (e.g., `sex`, `education_level`).

Understanding data types and null entries is essential before proceeding with analysis.

```
adult_df.shape
```

```
(32515, 16)
```

```
adult_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32515 entries, 0 to 32514
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32515 non-null  int64
1   workclass             32515 non-null  object
```

```

2   fnlwgt                32515 non-null  int64
3   education_num         32515 non-null  int64
4   marital_status       32515 non-null  object
5   relationship         32515 non-null  object
6   race                 32515 non-null  object
7   sex                  32515 non-null  object
8   capital_gain         32515 non-null  int64
9   capital_loss         32515 non-null  int64
10  hours_per_week       32515 non-null  int64
11  income               32515 non-null  object
12  education_level      32515 non-null  object
13  occupation_grouped   32515 non-null  object
14  native_region       32515 non-null  object
15  age_group            32515 non-null  object
dtypes: int64(6), object(10)
memory usage: 4.0+ MB

```

Summary Statistics: Numerical Variables

This summary provides a snapshot of key distribution characteristics. We see that:

- Age ranges from 17 to 90, with a mean of 38.6 years. It is slightly right-skewed (positively skewed). While the average age is approximately 38.6 years, an examination of the percentiles reveals that the majority of individuals are clustered in the younger to middle-age range, with fewer observations in the older age brackets. This skewed age distribution might suggest labor force participation is concentrated in specific age groups, which could reflect broader demographic or economic realities.
- Capital gains/losses are highly skewed, with most values at 0 (the 75th percentile is 0). This indicates that a small number of individuals report very large gains or losses, especially evident in the capital gain variable which reaches up to \$99,999. These variables act as proxies for wealth-related income that goes beyond regular wages or salaries. Individuals with non-zero values for capital gains or losses often represent a distinct socioeconomic subset of the population — typically more financially literate, or with access to investment assets. The stark inequality in their distributions mirrors real-world disparities in asset ownership and investment returns.
- The dataset has individuals working anywhere from 1 to 99 hours per week, with a median of 40. This aligns with the standard full-time work week in many countries (8 hours per day for 5 working days). The mean is slightly above that at 40.4 hours, suggesting a mild right skew, with a small subset of individuals working significantly longer hours. The mode is also 40, further reinforcing the prevalence of full-time work. A non-trivial number of individuals report working very few hours, possibly due to

part-time work, unemployment, or semi-retirement. On the other extreme, some report working more than 45 hours per week, which may indicate multiple jobs, weekend-work, self-employment, or informal labor, and could reflect socioeconomic necessity.

```
adult_df.describe()
```

	age	fnlwgt	education_num	capital_gain	capital_loss	hours_per_week
count	32515.000000	3.251500e+04	32515.000000	32515.000000	32515.000000	32515.000000
mean	38.590374	1.897912e+05	10.081593	1079.173428	87.427341	40.441089
std	13.638535	1.055766e+05	2.571943	7390.403187	403.231777	12.349830
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178300e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370475e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

Summary Statistics: Categorical Variables

workclass

The private sector dominates, employing ~69.7% of the population. The government sector (13.4%) and self-employment (11.2%) also make up substantial portions of the workforce. A small fraction is labeled as “unknown” (5.6%), which may correspond to missing or ambiguous data entries. Tiny proportions are voluntary (0.04%) or unemployed (0.02%), possibly underreported or underrepresented in the sample.

marital_status

Married individuals make up the largest group (46.1%), followed by those who are single (32.8%) and divorced or separated (18.1%). Widowed individuals represent a small minority (~3.1%).

relationship

The majority are labeled as “male spouse” (40.5%) or “single” (36.1%). Smaller categories include children (15.6%), female spouses (4.8%), and extended relatives (3.0%). The dominance of **male spouse** reflects the dataset’s gendered structure and may point to traditional family roles. The relative scarcity of “female spouse” roles suggests potential gender imbalances in how income-earning is reported within households.

race

The dataset is overwhelmingly composed of White individuals (~85.4%). Other racial groups include Black (9.6%), Asian or Pacific Islander (3.2%), American Indian or Eskimo (1.0%),

and Other (0.8%). The racial imbalance limits the generalizability of models trained on this data. Smaller racial groups may suffer from limited statistical power, affecting fairness and performance in predictive modeling.

sex

Males constitute 66.9% of the dataset, with females making up the remaining 33.1%. This male-skewed distribution could be due to sampling (e.g., primary earners in households), workforce participation patterns, or reporting biases.

education_level

Secondary-school graduates form the largest educational group (~32%), highlighting the central role of high school completion in the labor force. Tertiary education holders — those with university or equivalent degrees — account for nearly 25% of the population, representing a substantial segment with advanced qualifications. A notable 22.4% have attended some college without necessarily earning a degree, suggesting that partial post-secondary education is common, yet may not always translate into formal certification. The remaining 20% are distributed among those with only secondary education (9.4%), associate degrees (7.5%), primary school (3.5%), and a very small group with only preschool education (0.15%). It is evident that the education distribution is skewed toward mid- to high-level education, with relatively few individuals having only basic schooling. This reflects a dataset that largely captures working-age adults in formal labor, which may underrepresent the least-educated populations.

occupation_grouped

White-collar occupations are the most prevalent (~51%), followed by blue-collar, service, and unknown. Smaller categories include military, which is marginal. Essentially, slightly over half of individuals in the dataset work in professional, managerial, sales, clerical, or tech-support roles. This suggests the dataset is heavily weighted toward professional and administrative occupations. Nearly a third of the population works in manual labor or skilled trade positions (craft, transport, machine operation, farming, etc.). This indicates a significant segment engaged in physically intensive or technical labor.

native_region

The vast majority of individuals are from North America (~92.3%). Smaller proportions are from Central America, Asia, Europe, South America, and a generic Other category. The heavy concentration of North American individuals reflects the U.S. focus of the dataset.

age_group

The largest groups are 26–35 and 36–45, followed by 46–60. These three age groups represent about 73% of the dataset. Very few individuals are under 18 or above 75, consistent with the dataset’s focus on the working-age population.

```
adult_df.describe(include='object')
```

	workclass	marital_status	relationship	race	sex	income	education_level	occ
count	32515	32515	32515	32515	32515	32515	32515	32515
unique	7	4	6	5	2	2	7	5
top	private	married	husband	white	male	<=50k	secondary-school graduate	whi
freq	22652	14984	13178	27773	21760	24679	10485	165

```
adult_df['workclass'].value_counts(normalize=True)
```

```
workclass
private          0.696663
self-employed   0.112440
government       0.069414
local-gov        0.064370
unknown          0.056466
voluntary        0.000431
unemployed       0.000215
Name: proportion, dtype: float64
```

```
adult_df['marital_status'].value_counts(normalize=True)
```

```
marital_status
married          0.460833
single           0.327664
divorced or separated  0.180963
widowed          0.030540
Name: proportion, dtype: float64
```

```
adult_df['relationship'].value_counts(normalize=True)
```

```
relationship
husband          0.405290
not-in-family    0.254775
own-child        0.155590
unmarried        0.105951
wife             0.048224
other-relative   0.030171
Name: proportion, dtype: float64
```

```
adult_df['race'].value_counts(normalize=True)
```

```
race
white                0.854160
black                0.096017
asian or pacific islander  0.031924
american indian or eskimo  0.009565
other                0.008335
Name: proportion, dtype: float64
```

Income Distribution

Given that `income` is the target variable, most of the analysis hereafter will be based on it. We first of all examine the income distribution in the dataset.

```
adult_df_income = adult_df.groupby('income').size().reset_index(name = 'total')
adult_df_income
```

	income	total
0	<=50k	24679
1	>50k	7836

```
fig = px.pie(adult_df_income, names='income', values='total', title='Overall Income Distribution')
fig.update_layout(template="presentation", paper_bgcolor="rgba(0, 0, 0, 0)", plot_bgcolor = "#f5f5dc")
fig.show()
fig.write_image(os.path.join(results_dir, 'income_distribution_pie_chart.jpg'))
fig.write_image(os.path.join(results_dir, 'income_distribution_pie_chart.png'))
fig.write_html(os.path.join(results_dir, 'income_distribution_pie_chart.html'))
```

Overall Income Distribution



This pie chart visualizes the overall income split: 76% of individuals earn ≤50K, while 24% earn >50K. This means that nearly 3 out of 4 individuals fall into the lower income bracket (≤50K). This shows that there is a significant imbalance.

Income by Age Group

```
adult_df_income_age = adult_df.groupby(['age_group', 'income']).size().reset_index(name= 'total_by_age')
adult_df_income_age
```


	age_group	income	total_by_age
0	18-25	≤50k	5333
1	18-25	>50k	114
2	26-35	≤50k	6910
3	26-35	>50k	1591
4	36-45	≤50k	5232
5	36-45	>50k	2771
6	46-60	≤50k	4479
7	46-60	>50k	2809
8	61-75	≤50k	1580
9	61-75	>50k	511
10	76+	≤50k	200
11	76+	>50k	40
12	<18	≤50k	945

```
total_per_group = adult_df_income_age.groupby('age_group')['total_by_age'].transform('sum')
adult_df_income_age['percentage'] = (adult_df_income_age['total_by_age']/total_per_group)*100
adult_df_income_age
```

	age_group	income	total_by_age	percentage
0	18-25	≤50k	5333	97.907105
1	18-25	>50k	114	2.092895
2	26-35	≤50k	6910	81.284555
3	26-35	>50k	1591	18.715445
4	36-45	≤50k	5232	65.375484
5	36-45	>50k	2771	34.624516
6	46-60	≤50k	4479	61.457190
7	46-60	>50k	2809	38.542810

	age_group	income	total_by_age	percentage
8	61-75	<=50k	1580	75.561932
9	61-75	>50k	511	24.438068
10	76+	<=50k	200	83.333333
11	76+	>50k	40	16.666667
12	<18	<=50k	945	100.000000

```
fig = px.bar(
    adult_df_income_age,
    x = 'age_group',
    y = 'percentage',
    color = 'income',
    title = 'Income Distribution by Age Group (%)',
    barmode = 'group',
    color_discrete_sequence=['seagreen','lightgreen'],
    text = 'percentage'
)
fig.update_traces(texttemplate='%{text:.2f}%')
fig.update_layout(template = "presentation",
                    xaxis_title = 'native_region',
                    yaxis_title = 'percentage of population',
                    legend_title = dict(text='Income Level'),
                    paper_bgcolor="rgba(0, 0, 0, 0)",
                    plot_bgcolor = "rgba(0, 0, 0, 0)")
fig.show()
fig.write_image(os.path.join(results_dir, 'income_distribution_by_agegroup_bar_plot.jpg'))
fig.write_image(os.path.join(results_dir, 'income_distribution_by_agegroup_bar_plot.png'))
fig.write_html(os.path.join(results_dir, 'income_distribution_by_agegroup_bar_plot.html'))
```




The bar chart visualizes the income distribution across age groups, using percentages within

each group. There is an evident pattern in terms of income progression over the years with a gradual increase in terms of the number of people earning >50K starting from 0 amongst those aged 18 and below, peaking between 36 and 60 years, then declining after 60 years but not to zero.


All individuals under 18 earn <=50K, likely due to being students, minors, or ineligible for full-time employment. Extremely few young adults (2.1%) exceed 50K, as most are early in their careers, pursuing education, or in entry-level jobs. For the 26-35 age group, there's a noticeable improvement — roughly 1 in 5 individuals in this group earn >50K, reflecting early career progression and accumulation of qualifications/experience. A substantial income increase is seen in the 36-45 age group: over a third now earn >50K. This is typically considered prime earning age where individuals settle into stable, higher-paying positions. Highest proportion of >50K earners is seen amongst individuals aged between 46 and 60— nearly 4 in 10. This reflects career maturity, peak seniority levels, and accumulated experience. There's a drop-off in high incomes as many transition to retirement, part-time, or less demanding roles in the age group 61-75. Yet about 1 in 4 still earn >50K. Most in 76+ age group earn <=50K, likely due to retirement, pensions, or fixed incomes — but a small minority still earn higher incomes, possibly through continued work or investments.

```
themes = ["plotly", "plotly_white", "plotly_dark", "ggplot2", "seaborn", "simple_white", "pr
for theme in themes:
    fig.update_layout(template=theme)


fig.show()
```




Income Distribution by Age Group (%)









Income Distribution by Age Group (%)



Income Distribution by Age Group (%)







Income Distribution by native region

```
adult_df_income_reg = adult_df.groupby(['native_region', 'income']).size().reset_index(name=
adult_df_income_reg
```

	native_region	income	total_by_region
0	asia	<=50k	465
1	asia	>50k	206
2	central america	<=50k	466
3	central america	>50k	58
4	europa	<=50k	369
5	europa	>50k	152
6	north america	<=50k	22769
7	north america	>50k	7250
8	other	<=50k	436
9	other	>50k	146
10	south america	<=50k	174
11	south america	>50k	24

```
total_per_native_region = adult_df_income_reg.groupby('native_region')['total_by_region'].tr
adult_df_income_reg['percentage'] = (adult_df_income_reg['total_by_region']/total_per_native
adult_df_income_reg
```


	native_region	income	total_by_region	percentage
0	asia	<=50k	465	69.299553
1	asia	>50k	206	30.700447
2	central america	<=50k	466	88.931298
3	central america	>50k	58	11.068702
4	europa	<=50k	369	70.825336
5	europa	>50k	152	29.174664
6	north america	<=50k	22769	75.848629
7	north america	>50k	7250	24.151371
8	other	<=50k	436	74.914089
9	other	>50k	146	25.085911
10	south america	<=50k	174	87.878788
11	south america	>50k	24	12.121212

```
fig = px.bar(
    adult_df_income_reg,
    x = 'native_region',
    y = 'percentage',
    color = 'income',
    title = 'Income Distribution by native region (%)',
    barmode = 'group',
```

```

height= 540,
color_discrete_sequence=['seagreen', 'lightgreen'],
text = 'percentage'
)
fig.update_traces(texttemplate='%{text:.2f}%', textposition = 'outside')
fig.update_layout(template = "presentation",
                    axis_title = 'native_region',
                    yaxis_title = 'percentage of population',
                    legend_title = dict(text='Income Level'),
                    paper_bgcolor="rgba(0, 0, 0, 0)",
                    plot_bgcolor = "rgba(0, 0, 0, 0)")
fig.show()
fig.write_image(os.path.join(results_dir, 'income_distribution_by_nativeregion_bar_plot.jpg'))
fig.write_image(os.path.join(results_dir, 'income_distribution_by_nativeregion_bar_plot.png'))
fig.write_html(os.path.join(results_dir, 'income_distribution_by_nativeregion_bar_plot.html'))

```



Asia (30.7%) and Europe (29.2%) have the highest proportions of high-income earners. This suggests these immigrant groups might be better integrated into high-paying professional roles, or may represent a more skilled migrant profile in the dataset. Central America (11.1%) and South America (12.1%) have the lowest proportions of >50K earners. With 24.2% of North Americans earning >50K, this serves as a middle-ground baseline. Interestingly, both Asian and European groups outperform the native-born population proportionally in high-income brackets. The ‘Other’ group sits around 25.1%, close to North America’s rate. This likely reflects a diverse mix of regions not explicitly listed.

Income Distribution by Race


```
adult_df_income_race = adult_df.groupby(['race', 'income']).size().reset_index(name= 'total_by_race')
adult_df_income_race
```

	race	income	total_by_race
0	american indian or eskimo	<=50k	275
1	american indian or eskimo	>50k	36
2	asian or pacific islander	<=50k	762
3	asian or pacific islander	>50k	276
4	black	<=50k	2735
5	black	>50k	387
6	other	<=50k	246
7	other	>50k	25
8	white	<=50k	20661
9	white	>50k	7112

```
total_per_race = adult_df_income_race.groupby('race')['total_by_race'].transform('sum')
adult_df_income_race['percentage'] = (adult_df_income_race['total_by_race']/total_per_race)*100
adult_df_income_race
```


	race	income	total_by_race	percentage
0	american indian or eskimo	<=50k	275	88.424437
1	american indian or eskimo	>50k	36	11.575563
2	asian or pacific islander	<=50k	762	73.410405
3	asian or pacific islander	>50k	276	26.589595
4	black	<=50k	2735	87.604100
5	black	>50k	387	12.395900
6	other	<=50k	246	90.774908
7	other	>50k	25	9.225092
8	white	<=50k	20661	74.392395
9	white	>50k	7112	25.607605

```
fig = px.bar(
    adult_df_income_race,
    x = 'race',
    y = 'percentage',
    color = 'income',
    title = 'Income Distribution by race (%)',
    barmode = 'group',
```

```

height= 540,
width= 1200,
color_discrete_sequence=['seagreen','lightgreen'],
text = 'percentage'
)
fig.update_traces(texttemplate='%{text:.2f}%', textposition = 'outside')
fig.update_layout(template = "presentation",
                    xaxis_title = 'Race',
                    yaxis_title = 'percentage of population',
                    legend_title = dict(text='Income Level'),
                    paper_bgcolor="rgba(0, 0, 0, 0)",
                    plot_bgcolor = "rgba(0, 0, 0, 0)")
fig.show()

```



Asian or Pacific Islander (26.6%) and White (25.6%) populations have the highest proportions of >50K earners. Asians/Pacific Islanders marginally outperform Whites, a pattern often attributed to occupational concentration in high-paying sectors like technology and medicine. On the other hand, American Indian or Eskimo (11.6%), Black (12.4%), and Other (9.2%) groups show significantly lower rates of high-income earners. These figures reflect long-standing economic disparities rooted in historical exclusion, occupational segregation, and systemic inequality.

The stark differences in high-income proportions:

- **Between Whites and Blacks:** 25.6% vs 12.4% — slightly over double the proportion.

- **Between Asians and Others:** 26.6% vs 9.2% — nearly triple.

These disparities are consistent with well-documented wage gaps and underrepresentation of marginalized groups in higher-paying roles.

Income Distribution by Education Level and Occupation grouped

```
adult_df_income_edu_occ = adult_df.groupby(['education_level', 'occupation_grouped', 'income'])
adult_df_income_edu_occ
```

	education_level	occupation_grouped	income	total
33	secondary-school graduate	blue collar	<=50k	3977
62	tertiary	white collar	>50k	3545
61	tertiary	white collar	<=50k	3369
51	some college	white collar	<=50k	3004
41	secondary-school graduate	white collar	<=50k	2900
...
30	secondary	unknown	>50k	5
20	primary	unknown	>50k	4
13	preschool	white collar	<=50k	3
26	secondary	military	>50k	2
18	primary	service	>50k	1

```
adult_df_income_edu_occ['edu_occ'] = (adult_df_income_edu_occ['education_level'] + " | " + adult_df_income_edu_occ['occupation_grouped'])
adult_df_income_edu_occ
```

	education_level	occupation_grouped	income	total	edu_occ
33	secondary-school graduate	blue collar	<=50k	3977	secondary-school graduate blue collar
62	tertiary	white collar	>50k	3545	tertiary white collar
61	tertiary	white collar	<=50k	3369	tertiary white collar
51	some college	white collar	<=50k	3004	some college white collar
41	secondary-school graduate	white collar	<=50k	2900	secondary-school graduate white collar
...
30	secondary	unknown	>50k	5	secondary unknown
20	primary	unknown	>50k	4	primary unknown
13	preschool	white collar	<=50k	3	preschool white collar
26	secondary	military	>50k	2	secondary military

	education_level	occupation_grouped	income	total	edu_occ
18	primary	service	>50k	1	primary service


```

num = 15
adult_df_combos = adult_df_income_edu_occ.head(num)
fig = px.bar(
    adult_df_combos,
    x = 'total',
    y = 'edu_occ',
    color = 'income',
    orientation='h',
    title = f'Top {num} Education and Occupation Groups Combinations by Income group',
    #barmode = 'group',
    height= 500,
    width= 1100,
    color_discrete_sequence=['seagreen','lightgreen'],
    text = 'total'
)

fig.update_layout(template = "presentation",
                    xaxis_title = 'Number of Individuals',
                    yaxis_title = 'Education | Occupation group',
                    legend_title = dict(text='Income Level'),
                    margin = dict(l=450, r=50, t=50, b=50))
fig.update_traces(textposition = 'inside')
fig.show()
fig.write_image(os.path.join(results_dir, 'income_distribution_by_eduandocc_bar_plot.jpg'))
fig.write_image(os.path.join(results_dir, 'income_distribution_by_eduandocc_bar_plot.png'))
fig.write_html(os.path.join(results_dir, 'income_distribution_by_eduandocc_bar_plot.html'))

```

Top 15 Education and Occupation Groups Combinations by Income group



From the bar chart, we can pick out the largest groups per income-level. We see that secondary-school graduates working a blue collar job occupy the largest group in the dataset (3976). This reflects a common socio-economic profile: individuals with basic schooling in manual or technical trades predominantly earning lower incomes. The largest high-income group are tertiary-educated individuals in white collar roles. This highlights the strong earning advantage conferred by higher education and skilled jobs.

Some of the key patterns we can get from the dataset are:

- **Education matters, but isn't deterministic**

Tertiary education combined with white-collar work offers the highest income prospects. Yet a substantial number of tertiary-educated white-collar workers earn ≤50K, likely early career, part-time, or structural pay gaps.

- **Blue-collar and service work predominantly pay ≤50K, regardless of education.**

Even some college education doesn't guarantee high incomes in these sectors. Manual and service sector income is highly occupation-dependent (some skilled trades can break the 50K mark).

- **Some non-tertiary education groups do reach >50K**

Secondary-school graduates in blue-collar and white-collar work have decent representation among >50K earners. This reflects upward mobility possible through skilled trades, tenure, or niche roles.